

An Introduction to Data Wrangling with R (or, an Introduction to R with Data Wrangling)

Rod Alence
International Relations and e-Science
University of the Witwatersrand

Workshop
Introduction to Emerging Methodologies in Social Science Research (IEMeSSR)
10–14 May 2021

Aims of my sessions

My two sessions aim to introduce R as software for “data science” – understood to include:

1 Data “wrangling” (mainly this session)

- Importing (possibly “messy”) data into R – from text files, spreadsheets, or other statistical programs (SPSS, Stata, SAS, etc.); and
- “Tidying” data for analysis – getting the data into a rectangular (one-observation-per-row, one-variable-per-column) format.

2 Data analysis (mainly next session)

- Transforming data (mathematical calculations and recoding);
- Visualizing data (graphics); and
- Modeling data (statistics).

Approach of my sessions

“Learning by doing” (as much as possible)

- Start with “traditional” slides/lecture format;
- Shift to “live coding” in R/RStudio as soon as possible (technology permitting);
- Recommend options for offline self-study.

Programming within R

- Give “non-exclusive emphasis” to Base R over Tidyverse.
(If you have no idea what that means, don't worry!)

Expectations for my sessions (beyond attendance!)

What I do not expect

- Prior experience programming in R (a bonus if you have it!);
- Mathematical expertise beyond high-school algebra.

What I do expect

- Some prior experience using some statistical software (e.g., SPSS or Stata);
- Basic familiarity with descriptive statistics and statistical models (not beyond least-squares linear regression).

Outline

On learning R

Some R basics

- Data in R

- Functions in R

- Read data into R

Data “wrangling” with R: the UNDP Human Development Index

- Read the data into R

- “Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

Outline

On learning R

Some R basics

- Data in R

- Functions in R

- Read data into R

Data “wrangling” with R: the UNDP Human Development Index

- Read the data into R

- “Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

What is R?

- Free, open-source statistical software that runs on all major operating systems;
- Created in the mid-1990s at the University of Auckland, New Zealand, by Ross Ihaka and Robert Gentleman as an implementation of the S programming language;
- Now maintained by a volunteer Core Development Team, which releases an updated version about twice a year;
- New and updated add-on “packages” appear weekly – more than 17,000 now available;
- For more information: <http://www.r-project.org>

Why R?

- Is probably the most powerful software for statistical analysis;
- Has the best graphics capabilities;
- Its package system is “going viral” (in a good way);
- Is “free” – as intellectual property and in price.



[Home]

Download

[CRAN](#)

R Project

[About R](#)

[Logo](#)

[Contributors](#)

[What's New?](#)

[Reporting Bugs](#)

[Conferences](#)

[Search](#)

[Get Involved: Mailing Lists](#)

[Developer Pages](#)

[R Blog](#)

R Foundation

[Foundation](#)

[Board](#)

The R Project for Statistical Computing

Getting Started

R is a free software environment for statistical computing and graphics. It compiles and runs on a wide variety of UNIX platforms, Windows and MacOS. To [download R](#), please choose your preferred [CRAN mirror](#).

If you have questions about R like how to download and install the software, or what the license terms are, please read our [answers to frequently asked questions](#) before you send an email.

News

- [R version 4.1.0 \(Camp Pontanezen\) prerelease versions](#) will appear starting Saturday 2021-04-17. Final release is scheduled for Tuesday 2021-05-18.
- [R version 4.0.5 \(Shake and Throw\)](#) has been released on 2021-03-31.
- Thanks to the organisers of useR! 2020 for a successful online conference. Recorded tutorials and talks from the conference are available on the [R Consortium YouTube channel](#).
- [R version 3.6.3 \(Holding the Windsock\)](#) was released on 2020-02-29.
- You can support the R Foundation with a renewable subscription as a [supporting member](#)

News via Twitter

R package “task views”



<https://cran.r-project.org/web/views/>

CRAN Task Views

CRAN task views aim to provide some guidance which packages on CRAN are relevant for tasks related to automatically installed using the [ctv](#) package. The views are intended to have a sharp focus so that it is *not* meant to endorse the “best” packages for a given task.

- To automatically install the views, the [ctv](#) package needs to be installed, e.g., via
`install.packages("ctv")`
and then the views can be installed via `install.views` OR `update.views` (where the latter only installs those
`ctv::install.views("Econometrics")`
`ctv::update.views("Econometrics")`)
- The task views are maintained by volunteers. You can help them by suggesting packages that should be included in individual task view pages.
- For general concerns regarding task views contact the [ctv](#) package maintainer.

Topics

[Bayesian](#)

[ChemPhys](#)

[ClinicalTrials](#)

[Cluster](#)

[Databases](#)

[DifferentialEquations](#)

[Distributions](#)

[Econometrics](#)

[Environmetrics](#)

Bayesian Inference

Chemometrics and Computational Physics

Clinical Trial Design, Monitoring, and Analysis

Cluster Analysis & Finite Mixture Models

Databases with R

Differential Equations

Probability Distributions

Econometrics

Analysis of Ecological and Environmental Data

If statistics programs/languages were cars...



<https://twitter.com/statsepi/status/795574223439876100>

Does R have a “steep learning curve”?

The two most challenging things about R

- 1 It is entirely command (“expression”) based – you type commands, and R executes them (no “point-and-click” menus).
- 2 It allows multiple (unlimited) data “objects” in a session simultaneously.

But – these features are essential to R’s strengths

- No menu system could ever keep up with software as powerful and dynamic as R.
- Allowing multiple objects is essential to a programming language in which the output of nearly any command can be the input of another.

And... RStudio makes learning and executing R command syntax

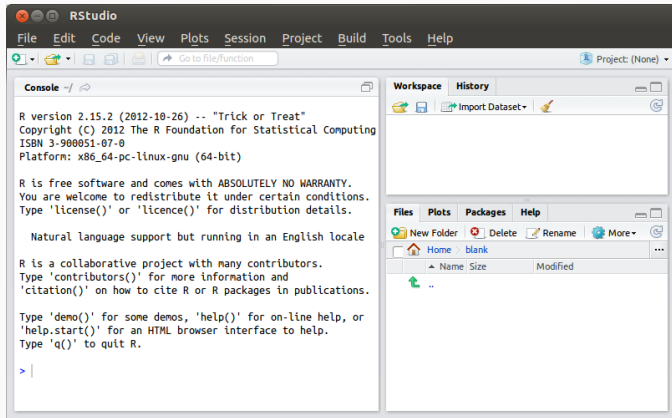
What is RStudio?

- An “integrated development environment” (IDE) for R (but not a “point-and-click” interface to R commands);
- Launched in 2011;
- Free and open source;
- Available for all major operating systems (Windows, MacOS, and Linux);
- For more information:
<http://www.rstudio.org>

RStudio at first start-up

Three windows

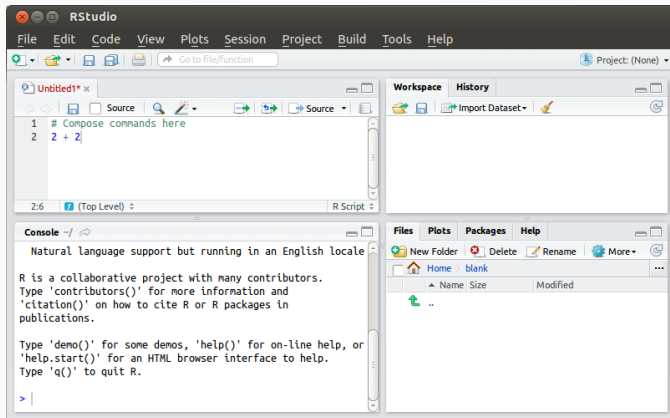
R console occupies full left side



RStudio with editor window open (the usual way)

Four windows

Left side split between editor (top) and R console (bottom)



Outline

On learning R

Some R basics

- Data in R

- Functions in R

- Read data into R

Data “wrangling” with R: the UNDP Human Development Index

- Read the data into R

- “Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

Outline

On learning R

Some R basics

- Data in R

- Functions in R

- Read data into R

Data “wrangling” with R: the UNDP Human Development Index

- Read the data into R

- “Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

Data structures: vectors

The most basic data structure in R is the **vector** – which is just a fancy word for “a list of things in a particular order.”

- Even a single number is a vector to R – it is a vector that happens to contain only one thing.
- When a vector holds data about units, like a column in a spreadsheet, a vector is synonymous with what is often called a “variable.”
- Each vector has two intrinsic structural attributes:

Length

How many things (elements) does it contain?

Mode

What kinds of things does it contain – e.g., numeric values, typographic characters?

Data structures: data frames

A **data frame** is a rectangular, spreadsheet-like data structure – typically organized with “observations” in rows and “variables” in columns.

Data frames

- **May** contain column vectors of **any class**; but
- **Must** contain column vectors of the **same length**.

The collection of packages known as the “Tidyverse” often use a special type of data frame called the **tibble**, which

- Has the same basic structure as the “traditional” data frame, with a few distinct features; and
- Can easily be converted to the “traditional” data frame.

Assigning names to data

Any non-trivial “data” is **assigned** a name for further use, using the “backward-arrow” operator.

For example, we can “stick together” some numbers as a vector using `c` (for combine) and assign it a name, like `some_numbers`:

```
some_numbers <- c(1, 2, 3, 4, 5)
```

And we can do the same with some letters (in quotation marks):

```
some_letters <- c("e", "d", "c", "b", "a")
```

Assignment is “silent” – but we can check the objects’ contents by typing its name and pressing return.

```
some_numbers  
# [1] 1 2 3 4 5
```

```
some_letters  
# [1] "e" "d" "c" "b" "a"
```

Combining vectors in a data frame

Because the two vectors assigned in the previous slide are the same length, we can stick them together side-by-side in a data frame using `data.frame`.

```
boring_df <- data.frame(some_numbers, some_letters)
```

And to view the data frame, enter its name.

```
boring_df
#   some_numbers some_letters
# 1             1           e
# 2             2           d
# 3             3           c
# 4             4           b
# 5             5           a
```

Notes on naming

R's rules about names

- May contain lower-case and upper-case letters, numbers, dots (.), and underscores (_).
- May not start with numbers – and they should almost always start with letters.
- Are case-sensitive – lower-case and upper-case versions of the same letter are treated as entirely different characters.
- Overwrite any existing object with the same name.

Common sense about names

- Should be concise (to avoid too much typing).
- Should be Informative (to clarify content).

Numeric indexing

Elements of data structures can be accessed by position using numeric square-bracket indexes.

Vectors

```
some_letters  
# [1] "e" "d" "c" "b" "a"  
  
some_letters[4] # get the fourth element  
# [1] "b"
```

Data frames

```
boring_df  
#   some_numbers some_letters  
# 1             1           e  
# 2             2           d  
# 3             3           c  
# 4             4           b  
# 5             5           a  
  
boring_df[3, 2] # row-by-col (third row, second col)  
# [1] "c"
```

Indexing columns (variables) in data frames

Three common ways to select a (column) variable in a data frame:

1 By numeric position

```
boring_df[ , 2] # get the second col (all rows)  
# [1] "e" "d" "c" "b" "a"
```

2 By column name

```
boring_df[ , "some_letters"] # get the col called "some_letters"  
# [1] "e" "d" "c" "b" "a"
```

3 Dollar-sign (list) notation

```
boring_df$some_letters  
# [1] "e" "d" "c" "b" "a"
```


Indexing rows (observations) in data frames

Two common ways to select rows in a data frame:

1 By numeric position

```
boring_df[c(1, 3), ] # get the first and third rows (all cols)
#   some_numbers some_letters
# 1             1           e
# 3             3           c
```

2 By logical expression

```
## Get rows in which the logical expression holds
boring_df[boring_df$some_numbers > 3, ]
#   some_numbers some_letters
# 4             4           b
# 5             5           a
```

Outline

On learning R

Some R basics

Data in R

Functions in R

Read data into R

Data “wrangling” with R: the UNDP Human Development Index

Read the data into R

“Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

Functions in R

Functions are what “do things” in R – if data objects are like nouns, functions are the verbs.

Function syntax

To use a function:

- 1 Type its name (exactly, remember case-sensitivity),
- 2 Followed immediately by parentheses (curved brackets),
- 3 Insert any inputs (“arguments”) inside the parentheses, separated by commas.

Often the reason for using a function is to generate output which is immediately assigned to an object.

An example using functions

Which two functions are used here?

```
marks <- c(78, 56, 91, 88, NA, 62, 67) # one student was absent
class_ave <- mean(marks, na.rm=TRUE)
class_ave
# [1] 73.66667
```

Each function has a help page, which explains what the function does and what inputs it takes.

Typing a question mark followed by a function name calls up the help page. To find out what the `na.rm=TRUE` is about, try entering `?mean` in the R console.

Outline

On learning R

Some R basics

Data in R

Functions in R

Read data into R

Data “wrangling” with R: the UNDP Human Development Index

Read the data into R

“Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

Reading data into R

R has functions for reading in data in various formats, for example:

- **Comma- or tab-delimited text:** `read.csv` and `read.delim` in Base R, or `read_csv` and `read_tsv` in the `readr` package
- **Spreadsheets (.xls, .xlsx):** `read_excel` in the `readxl` package;
- **Stata (.dta):** `read.dta` (Stata 5-12) in the `foreign` package, `read.dta13` (Stata 13 onwards) in the `readstata13` package, or `read_dta` (all versions) in the `haven` package;
- **SPSS (.sav):** `read.spss` in the `foreign` package, or `read_spss` in the `haven` package.

Outline

On learning R

Some R basics

- Data in R

- Functions in R

- Read data into R

Data “wrangling” with R: the UNDP Human Development Index

- Read the data into R

- “Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

Setting up

Install two packages

The installation only needs to be run once – I use the console.

```
install.packages(c("readxl", "tidyverse"))
```

Download the data spreadsheet

The download only needs to be run once (if the data set is static).

```
## Break up the url for convenience, because it is long
site_url <- "http://hdr.undp.org/"
path_url <- "sites/default/files/"
fn_url    <- "2020_statistical_annex_table_1.xlsx"
## Paste the parts together
link_url <- paste0(site_url, path_url, fn_url)

## Download using the full url
download.file(url=link_url,
              destfile="hdi2020.xlsx")
```


Outline

On learning R

Some R basics

- Data in R

- Functions in R

- Read data into R

Data “wrangling” with R: the UNDP Human Development Index

- Read the data into R

- “Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

Read in the spreadsheet (and have a quick look)

```
library(readxl)
HDI <- read_excel("hdi2020.xlsx",
                  range = "B8:K200", # cells to read
                  col_names = FALSE, # column names not read
                  na = c("", " ")) # missing value strings

dim(HDI)
# [1] 193 10

head(HDI)
# # A tibble: 6 x 10
#   ... 1      ... 2    ... 3    ... 4    ... 5    ... 6    ... 7    ... 8    ... 9
#   <chr>    <dbl> <lg1> <dbl> <chr> <dbl> <chr> <dbl> <chr>
# 1 VERY HIG~ NA      NA      NA <NA>   NA <NA>   NA <NA>
# 2 Norway    0.957 NA      82.4 <NA>   18.1 b    12.9 <NA>
# 3 Ireland   0.955 NA      82.3 <NA>   18.7 b    12.7 <NA>
# 4 Switzerl~ 0.955 NA      83.8 <NA>   16.3 <NA>  13.4 <NA>
# 5 Hong Kon~ 0.949 NA      84.9 <NA>   16.9 <NA>  12.3 <NA>
# 6 Iceland   0.949 NA      83.0 <NA>   19.1 b    12.8 c
# # ... with 1 more variable: ... 10 <dbl>
```

The structure of the data frame

The data frame is still a bit “messy.”

```
str(HDI)
# tibble[,10] [193 x 10] (S3: tbl_df/tbl/data.frame)
# $ ...1 : chr [1:193] "VERY HIGH HUMAN DEVELOPMENT" "Norway" "Ireland"
# $ ...2 : num [1:193] NA 0.957 0.955 0.955 0.949 0.949 0.947 0.945 0.
# $ ...3 : logi [1:193] NA NA NA NA NA NA NA ...
# $ ...4 : num [1:193] NA 82.4 82.3 83.8 84.9 ...
# $ ...5 : chr [1:193] NA NA NA NA ...
# $ ...6 : num [1:193] NA 18.1 18.7 16.3 16.9 ...
# $ ...7 : chr [1:193] NA "b" "b" NA ...
# $ ...8 : num [1:193] NA 12.9 12.7 13.4 12.3 ...
# $ ...9 : chr [1:193] NA NA NA NA ...
# $ ...10: num [1:193] NA 66494 68371 69394 62985 ...
```

Outline

On learning R

Some R basics

- Data in R

- Functions in R

- Read data into R

Data “wrangling” with R: the UNDP Human Development Index

- Read the data into R

- “Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

Remove unneeded columns

Which columns are not needed?

```
head(HDI)
# # A tibble: 6 x 10
#   ... 1      ... 2 ... 3    ... 4 ... 5    ... 6 ... 7    ... 8 ... 9
#   <chr>    <dbl> <lgl> <dbl> <chr> <dbl> <chr> <dbl> <chr>
# 1 VERY HIG~ NA      NA      NA <NA>   NA <NA>   NA <NA>
# 2 Norway    0.957 NA      82.4 <NA>   18.1 b     12.9 <NA>
# 3 Ireland   0.955 NA      82.3 <NA>   18.7 b     12.7 <NA>
# 4 Switzerl~ 0.955 NA      83.8 <NA>   16.3 <NA>   13.4 <NA>
# 5 Hong Kon~ 0.949 NA      84.9 <NA>   16.9 <NA>   12.3 <NA>
# 6 Iceland   0.949 NA      83.0 <NA>   19.1 b     12.8 c
# # ... with 1 more variable: ... 10 <dbl>
```

Remove the “skinny” footnote columns.

```
## Use negative column indexes to remove columns
HDI <- HDI[ , -c(3, 5, 7, 9)] # negative column index to remove

## Tidyverse alternative (dplyr package) (NOT RUN)
## HDI <- dplyr::select(HDI, -c(3, 5, 7, 9))
```

Add meaningful **column names**

Add meaningful (informative and concise) names by “assigning into” the data frame’s `colnames`:

```
## Old column names
colnames(HDI)
# [1] " ... 1" " ... 2" " ... 4" " ... 6" " ... 8" " ... 10"

## Assign new column names
colnames(HDI) <- c("country", "hdi", "life_exp",
                  "school_exp", "school_mean", "gni_pc")

head(HDI)
# # A tibble: 6 x 6
#   country      hdi life_exp school_exp school_mean gni_pc
#   <chr>      <dbl>   <dbl>      <dbl>      <dbl>   <dbl>
# 1 VERY HIGH H~ NA      NA      NA      NA      NA
# 2 Norway      0.957    82.4    18.1    12.9 66494.
# 3 Ireland     0.955    82.3    18.7    12.7 68371.
# 4 Switzerland 0.955    83.8    16.3    13.4 69394.
# 5 Hong Kong, ~ 0.949    84.9    16.9    12.3 62985.
# 6 Iceland     0.949    83.0    19.1    12.8 54682.
```

Remove unneeded rows

Which rows are not needed?

```
head(HDI, n=2)
# # A tibble: 2 x 6
#   country      hdi life_exp school_exp school_mean gni_pc
#   <chr>      <dbl>   <dbl>     <dbl>     <dbl>   <dbl>
# 1 VERY HIGH H~ NA       NA       NA       NA       NA
# 2 Norway     0.957    82.4     18.1     12.9    66494.

HDI[HDI$country == "MEDIUM HUMAN DEVELOPMENT", ]
# # A tibble: 1 x 6
#   country      hdi life_exp school_exp school_mean gni_pc
#   <chr>      <dbl>   <dbl>     <dbl>     <dbl>   <dbl>
# 1 MEDIUM HUMAN~ NA       NA       NA       NA       NA
```

Remove rows with no numeric data (e.g., in the hdi column).

```
HDI <- HDI[! is.na(HDI$hdi), ]

## Tidyverse alternative (NOT RUN)
## HDI <- dplyr::filter(HDI, ! is.na(HDI$hdi))
```

Check the data structure

Things to check:

- Dimensions (rows by columns);
- Column names;
- Object “classes” (e.g., character vs. numeric).

```
str(HDI)
```

```
# tibble[,6] [189 x 6] (S3: tbl_df/tbl/data.frame)
```

```
# $ country      : chr [1:189] "Norway" "Ireland" "Switzerland" "Hong Ko
```

```
# $ hdi          : num [1:189] 0.957 0.955 0.955 0.949 0.949 0.947 0.945
```

```
# $ life_exp     : num [1:189] 82.4 82.3 83.8 84.9 83 ...
```

```
# $ school_exp  : num [1:189] 18.1 18.7 16.3 16.9 19.1 ...
```

```
# $ school_mean: num [1:189] 12.9 12.7 13.4 12.3 12.8 ...
```

```
# $ gni_pc      : num [1:189] 66494 68371 69394 62985 54682 ...
```


Check the data **summary**


Things to check:

- Descriptive statistics;
- Missing values (NA frequencies are reported for each variable).

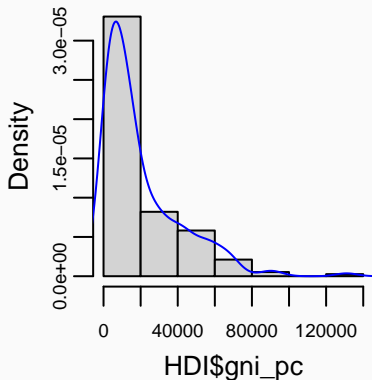
```
summary(HDI)
#      country                hdi                life_exp
# Length:189             Min.   :0.3940             Min.   :53.28
# Class :character       1st Qu.:0.6020             1st Qu.:67.44
# Mode  :character       Median :0.7400             Median :74.05
#                               Mean  :0.7224             Mean  :72.71
#                               3rd Qu.:0.8290             3rd Qu.:77.91
#                               Max.   :0.9570             Max.   :84.86
#      school_exp      school_mean      gni_pc
# Min.   : 5.005      Min.   : 1.644      Min.   : 753.9
# 1st Qu.:11.431      1st Qu.: 6.437      1st Qu.: 4910.2
# Median :13.188      Median : 9.032      Median : 12707.4
# Mean   :13.325      Mean   : 8.728      Mean   : 20219.7
# 3rd Qu.:15.227      3rd Qu.:11.326      3rd Qu.: 29497.2
# Max.   :21.954      Max.   :14.152      Max.   :131031.6
```

Run a few **hist**ograms?

```
hist(HDI$hdi,  
     main="")
```



```
hist(HDI$gni_pc,  
     freq=FALSE,  
     main="")  
lines(density(HDI$gni_pc),  
      col="blue")
```



Outline

On learning R

Some R basics

- Data in R

- Functions in R

- Read data into R

Data “wrangling” with R: the UNDP Human Development Index

- Read the data into R

- “Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

What is `swirl`

`swirl`

An R package that provides the infrastructure to run interactive self-study lessons, right in the R console.

(For more information: <https://swirlstats.com/>.)

“Introduction to R Programming”

A foundational `swirl` course consisting of 15 short lessons (work out your own pace, but most take about 15–20 minutes each).

Install and run swirl courses

Install (once-off)

```
## Install the swirl package -- the "infrastructure"  
install.packages("swirl")  
  
## Install the R programming course -- the content  
swirl::install_course("R Programming")
```

Run the course

```
## "Load" (attach) the swirl package  
library(swirl)  
  
## Run swirl  
swirl() # follow the prompts to choose a course, lessons
```

(Mostly do the courses in order, except you may want to skip the second one on “Workspace and Files” and come back to it later.)

Outline

On learning R

Some R basics

- Data in R

- Functions in R

- Read data into R

Data “wrangling” with R: the UNDP Human Development Index

- Read the data into R

- “Tidy” the data

Self-study in R/RStudio with `swirl`

Wrapping up

Wrapping up

Learning R

- Push ahead with the `swirl` lessons on “R Programming”;
- For further self-study, with a Tidyverse focus, try:
Hadley Wickham and Garrett Grolemund, *R for Data Science* (O'Reilly Media, 2017) – available for free online at <https://r4ds.had.co.nz/>.

For tomorrow

- Data *analysis* in R.