

# An Introduction to Data Wrangling and Analysis with R

---

Rod Alence  
International Relations and e-Science  
University of the Witwatersrand

Workshop  
Introduction to Emerging Methodologies in Social Science Research (IEMeSSR)  
10–14 May 2021

# Aims of my sessions

My two sessions aim to introduce R as software for “data science” – understood to include:

## 1 Data “wrangling”

- Importing (possibly “messy”) data into R – from text files, spreadsheets, or other statistical programs (SPSS, Stata, SAS, etc.); and
- “Tidying” data for analysis – getting the data into a rectangular (one-observation-per-row, one-variable-per-column) format.

## 2 Data analysis

- Transforming data (mathematical calculations and recoding);
- Visualizing data (graphics); and
- Modeling data (statistics).

# Approach of my sessions

## “Learning by doing” (as much as possible)

- Start with “traditional” slides/lecture format;
- Shift to “live coding” in R/RStudio as soon as possible (technology permitting);
- Recommend options for offline self-study.

## Programming within R

- Give “non-exclusive emphasis” to Base R over Tidyverse.  
(If you have no idea what that means, don't worry!)

## Expectations for my sessions (beyond attendance!)

### What I do not expect

- Prior experience programming in R (a bonus if you have it!);
- Mathematical expertise beyond high-school algebra.

### What I do expect

- Some prior experience using some statistical software (e.g., SPSS or Stata);
- Basic familiarity with descriptive statistics and statistical models (not beyond least-squares linear regression).

# Outline

On learning R

Some R basics

Data “wrangling” with R: the UNDP Human Development Index

Model and visualize the data

Self-study in R/RStudio with `swirl`

Wrapping up

# Outline

On learning R

Some R basics

Data “wrangling” with R: the UNDP Human Development Index

Model and visualize the data

Self-study in R/RStudio with `swirl`

Wrapping up


# What is R?

- Free, open-source statistical software that runs on all major operating systems;
- Created in the mid-1990s at the University of Auckland, New Zealand, by Ross Ihaka and Robert Gentleman as an implementation of the S programming language;
- Now maintained by a volunteer Core Development Team, which releases an updated version about twice a year;
- New and updated add-on “packages” appear weekly – more than 17,000 now available;
- For more information: <http://www.r-project.org>

## Why R?

- Is probably the most powerful software for statistical analysis;
- Has the best graphics capabilities;
- Its package system is “going viral” (in a good way);
- Is “free” – as intellectual property and in price.





[Home]

**Download**

CRAN

**R Project**

About R

Logo

Contributors

What's New?

Reporting Bugs

Conferences

Search

Get Involved: Mailing Lists

Developer Pages

R Blog

**R Foundation**

Foundation

Board

## The R Project for Statistical Computing

### Getting Started

R is a free software environment for statistical computing and graphics. It compiles and runs on a wide variety of UNIX platforms, Windows and MacOS. To [download R](#), please choose your preferred [CRAN mirror](#).

If you have questions about R like how to download and install the software, or what the license terms are, please read our [answers to frequently asked questions](#) before you send an email.

### News

- [R version 4.1.0 \(Camp Pontanezen\) prerelease versions](#) will appear starting Saturday 2021-04-17. Final release is scheduled for Tuesday 2021-05-18.
- [R version 4.0.5 \(Shake and Throw\)](#) has been released on 2021-03-31.
- Thanks to the organisers of useR! 2020 for a successful online conference. Recorded tutorials and talks from the conference are available on the [R Consortium YouTube channel](#).
- [R version 3.6.3 \(Holding the Windsock\)](#) was released on 2020-02-29.
- You can support the R Foundation with a renewable subscription as a [supporting member](#)

### News via Twitter

# R package “task views”



<https://cran.r-project.org/web/views/>

## CRAN Task Views

CRAN task views aim to provide some guidance which packages on CRAN are relevant for tasks related to automatically installed using the [ctv](#) package. The views are intended to have a sharp focus so that it is *not* meant to endorse the “best” packages for a given task.

- To automatically install the views, the [ctv](#) package needs to be installed, e.g., via  
`install.packages("ctv")`  
and then the views can be installed via `install.views` OR `update.views` (where the latter only installs those  
`ctv::install.views("Econometrics")`  
`ctv::update.views("Econometrics")`)
- The task views are maintained by volunteers. You can help them by suggesting packages that should be included in individual task view pages.
- For general concerns regarding task views contact the [ctv](#) package maintainer.

### Topics

[Bayesian](#)

[ChemPhys](#)

[ClinicalTrials](#)

[Cluster](#)

[Databases](#)

[DifferentialEquations](#)

[Distributions](#)

[Econometrics](#)

[Environmetrics](#)

Bayesian Inference

Chemometrics and Computational Physics

Clinical Trial Design, Monitoring, and Analysis

Cluster Analysis & Finite Mixture Models

Databases with R

Differential Equations

Probability Distributions

Econometrics

Analysis of Ecological and Environmental Data

# If statistics programs/languages were cars...



<https://twitter.com/statsepi/status/795574223439876100>

# Does R have a “steep learning curve”?

## The two most challenging things about R

- 1 It is entirely command (“expression”) based – you type commands, and R executes them (no “point-and-click” menus).
- 2 It allows multiple (unlimited) data “objects” in a session simultaneously.

## But – these features are essential to R’s strengths

- No menu system could ever keep up with software as powerful and dynamic as R.
- Allowing multiple objects is essential to a programming language in which the output of nearly any command can be the input of another.

And... RStudio makes learning and executing R command syntax

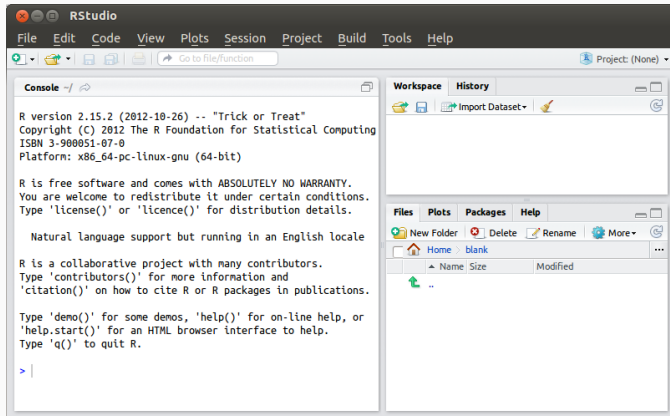
# What is RStudio?

- An “integrated development environment” (IDE) for R (but not a “point-and-click” interface to R commands);
- Launched in 2011;
- Free and open source;
- Available for all major operating systems (Windows, MacOS, and Linux);
- For more information:  
<http://www.rstudio.org>

# RStudio at first start-up

## Three windows

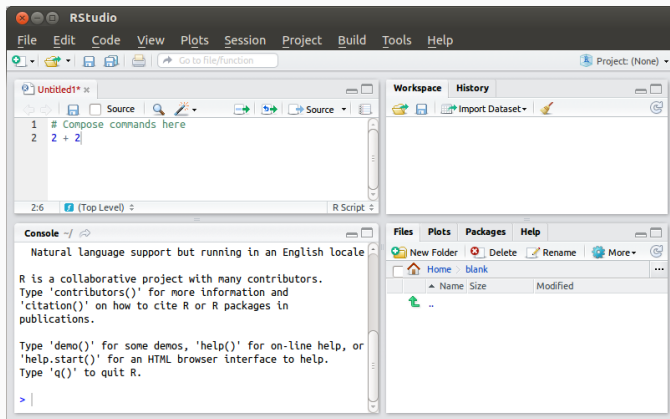
R console occupies full left side



# RStudio with editor window open (the usual way)

## Four windows

Left side split between editor (top) and R console (bottom)



# Outline

On learning R

**Some R basics**

Data “wrangling” with R: the UNDP Human Development Index

Model and visualize the data

Self-study in R/RStudio with `swirl`

Wrapping up



## Data structures: vectors

The most basic data structure in R is the **vector** – which is just a fancy word for “a list of things in a particular order.”

- Even a single number is a vector to R – it is a vector that happens to contain only one thing.
- When a vector holds data about units, like a column in a spreadsheet, a vector is synonymous with what is often called a “variable.”
- Each vector has two intrinsic structural attributes:

### Length

How many things (elements) does it contain?

### Mode

What kinds of things does it contain – e.g., numeric values, typographic characters?

## Data structures: data frames

A **data frame** is a rectangular, spreadsheet-like data structure – typically organized with “observations” in rows and “variables” in columns.

### Data frames

- **May** contain column vectors of **any class**; but
- **Must** contain column vectors of the **same length**.

The collection of packages known as the “Tidyverse” often use a special type of data frame called the **tibble**, which

- Has the same basic structure as the “traditional” data frame, with a few distinct features; and
- Can easily be converted to the “traditional” data frame.

## Assigning names to data

Any non-trivial “data” is **assigned** a name for further use, using the “backward-arrow” operator.

For example, we can “stick together” some numbers as a vector using `c` (for combine) and assign it a name, like `some_numbers`:

```
some_numbers <- c(1, 2, 3, 4, 5)
```

And we can do the same with some letters (in quotation marks):

```
some_letters <- c("e", "d", "c", "b", "a")
```

Assignment is “silent” – but we can check the objects’ contents by typing its name and pressing return.

```
some_numbers  
# [1] 1 2 3 4 5
```

```
some_letters  
# [1] "e" "d" "c" "b" "a"
```

## Combining vectors in a data frame

Because the two vectors assigned in the previous slide are the same length, we can stick them together side-by-side in a data frame using `data.frame`.

```
boring_df <- data.frame(some_numbers, some_letters)
```

And to view the data frame, enter its name.

```
boring_df
#   some_numbers some_letters
# 1             1           e
# 2             2           d
# 3             3           c
# 4             4           b
# 5             5           a
```

# Notes on naming

## R's rules about names

- May contain lower-case and upper-case letters, numbers, dots (.), and underscores (\_).
- May not start with numbers – and they should almost always start with letters.
- Are case-sensitive – lower-case and upper-case versions of the same letter are treated as entirely different characters.
- Overwrite any existing object with the same name.

## Common sense about names

- Should be concise (to avoid too much typing).
- Should be Informative (to clarify content).

# Numeric indexing

Elements of data structures can be accessed by position using numeric square-bracket indexes.

## Vectors

```
some_letters  
# [1] "e" "d" "c" "b" "a"  
  
some_letters[4] # get the fourth element  
# [1] "b"
```

## Data frames

```
boring_df  
#   some_numbers some_letters  
# 1             1           e  
# 2             2           d  
# 3             3           c  
# 4             4           b  
# 5             5           a  
  
boring_df[3, 2] # row-by-col (third row, second col)  
# [1] "c"
```

# Indexing columns (variables) in data frames

Three common ways to select a column (variable) in a data frame:

## 1 By numeric position

```
boring_df[ , 2] # get the second col (all rows)  
# [1] "e" "d" "c" "b" "a"
```

## 2 By column name

```
boring_df[ , "some_letters"] # get the col called "some_letters"  
# [1] "e" "d" "c" "b" "a"
```

## 3 Dollar-sign (list) notation

```
boring_df$some_letters  
# [1] "e" "d" "c" "b" "a"
```

# Indexing rows (observations) in data frames

Two common ways to select rows in a data frame:

## 1 By numeric position

```
boring_df[c(1, 3), ] # get the first and third rows (all cols)
#   some_numbers some_letters
# 1             1           e
# 3             3           c
```

## 2 By logical expression

```
## Get rows in which the logical expression holds
boring_df[boring_df$some_numbers > 3, ]
#   some_numbers some_letters
# 4             4           b
# 5             5           a
```



# Functions in R

Functions are what “do things” in R – if data objects are like nouns, functions are the verbs.

## Function syntax

To use a function:

- 1 Type its name (exactly, remember case-sensitivity),
- 2 Followed immediately by parentheses (curved brackets),
- 3 Insert any inputs (“arguments”) inside the parentheses, separated by commas.

Often the reason for using a function is to generate output which is immediately assigned to an object.

## An example using functions

Which two functions are used here?

```
marks <- c(78, 56, 91, 88, NA, 62, 67) # one student was absent
class_ave <- mean(marks, na.rm=TRUE)
class_ave
# [1] 73.66667
```

Each function has a help page, which explains what the function does and what inputs it takes.

Typing a question mark followed by a function name calls up the help page. To find out what the `na.rm=TRUE` is about, try entering `?mean` in the R console.

## Reading data into R

R has functions for reading in data in various formats, for example:

- **Comma- or tab-delimited text:** `read.csv` and `read.delim` in Base R, or `read_csv` and `read_tsv` in the `readr` package
- **Spreadsheets (.xls, .xlsx):** `read_excel` in the `readxl` package;
- **Stata (.dta):** `read.dta` (Stata 5-12) in the `foreign` package, `read.dta13` (Stata 13 onwards) in the `readstata13` package, or `read_dta` (all versions) in the `haven` package;
- **SPSS (.sav):** `read.spss` in the `foreign` package, or `read_spss` in the `haven` package.

# Outline

On learning R

Some R basics

Data “wrangling” with R: the UNDP Human Development Index

Model and visualize the data

Self-study in R/RStudio with `swirl`

Wrapping up

# Setting up

## Install two packages

The installation only needs to be run once – I use the console.

```
install.packages(c("readxl", "tidyverse"))
```

## Download the data spreadsheet

The download only needs to be run once (if the data set is static).

```
## Break up the url for convenience, because it is long
site_url <- "http://hdr.undp.org/"
path_url <- "sites/default/files/"
fn_url    <- "2020_statistical_annex_table_1.xlsx"

## Paste the parts together
link_url <- paste0(site_url, path_url, fn_url)

## Download using the full url
download.file(url=link_url,
              method="curl",
              destfile="hdi2020.xlsx")
```

# Read in the spreadsheet (and have a quick look)

```
library(readxl)
HDI <- read_excel("hdi2020.xlsx",
                  range = "B8:K200", # cells to read
                  col_names = FALSE, # column names not read
                  na = c("", " ")) # missing value strings

dim(HDI)
# [1] 193 10

head(HDI)
# # A tibble: 6 x 10
#   ... 1      ... 2    ... 3    ... 4    ... 5    ... 6    ... 7    ... 8    ... 9
#   <chr>    <dbl> <lg1> <dbl> <chr> <dbl> <chr> <dbl> <chr>
# 1 VERY HIG~ NA      NA      NA <NA>  NA <NA>  NA <NA>
# 2 Norway    0.957 NA      82.4 <NA>  18.1 b    12.9 <NA>
# 3 Ireland   0.955 NA      82.3 <NA>  18.7 b    12.7 <NA>
# 4 Switzerl~ 0.955 NA      83.8 <NA>  16.3 <NA>  13.4 <NA>
# 5 Hong Kon~ 0.949 NA      84.9 <NA>  16.9 <NA>  12.3 <NA>
# 6 Iceland   0.949 NA      83.0 <NA>  19.1 b    12.8 c
# # ... with 1 more variable: ... 10 <dbl>
```

# The structure of the data frame

The data frame is still a bit “messy.”

```
str(HDI)
# tibble[,10] [193 x 10] (S3: tbl_df/tbl/data.frame)
# $ ...1 : chr [1:193] "VERY HIGH HUMAN DEVELOPMENT" "Norway" "Ireland"
# $ ...2 : num [1:193] NA 0.957 0.955 0.955 0.949 0.949 0.947 0.945 0.
# $ ...3 : logi [1:193] NA NA NA NA NA NA NA ...
# $ ...4 : num [1:193] NA 82.4 82.3 83.8 84.9 ...
# $ ...5 : chr [1:193] NA NA NA NA ...
# $ ...6 : num [1:193] NA 18.1 18.7 16.3 16.9 ...
# $ ...7 : chr [1:193] NA "b" "b" NA ...
# $ ...8 : num [1:193] NA 12.9 12.7 13.4 12.3 ...
# $ ...9 : chr [1:193] NA NA NA NA ...
# $ ...10: num [1:193] NA 66494 68371 69394 62985 ...
```

# Remove unneeded columns

## Which columns are not needed?

```
head(HDI)
# # A tibble: 6 x 10
#   ... 1      ... 2 ... 3    ... 4 ... 5    ... 6 ... 7    ... 8 ... 9
#   <chr>    <dbl> <lgl> <dbl> <chr> <dbl> <chr> <dbl> <chr>
# 1 VERY HIG~ NA      NA      NA <NA>   NA <NA>   NA <NA>
# 2 Norway    0.957 NA      82.4 <NA>   18.1 b     12.9 <NA>
# 3 Ireland   0.955 NA      82.3 <NA>   18.7 b     12.7 <NA>
# 4 Switzerl~ 0.955 NA      83.8 <NA>   16.3 <NA>   13.4 <NA>
# 5 Hong Kon~ 0.949 NA      84.9 <NA>   16.9 <NA>   12.3 <NA>
# 6 Iceland   0.949 NA      83.0 <NA>   19.1 b     12.8 c
# # ... with 1 more variable: ... 10 <dbl>
```

## Remove the “skinny” footnote columns.

```
## Use negative column indexes to remove columns
HDI <- HDI[ , -c(3, 5, 7, 9)] # negative column index to remove

## Tidyverse alternative (dplyr package) (NOT RUN)
## HDI <- dplyr::select(HDI, -c(3, 5, 7, 9))
```



## Add meaningful **column names**

Add meaningful (informative and concise) names by “assigning into” the data frame’s `colnames`:

```
## Old column names
colnames(HDI)
# [1] " ... 1"  " ... 2"  " ... 4"  " ... 6"  " ... 8"  " ... 10"

## Assign new column names
colnames(HDI) <- c("country", "hdi", "life_exp",
                  "school_exp", "school_mean", "gni_pc")

head(HDI)
# # A tibble: 6 x 6
#   country      hdi life_exp school_exp school_mean gni_pc
#   <chr>      <dbl>   <dbl>      <dbl>      <dbl>   <dbl>
# 1 VERY HIGH H~ NA      NA      NA      NA      NA
# 2 Norway      0.957    82.4    18.1    12.9 66494.
# 3 Ireland     0.955    82.3    18.7    12.7 68371.
# 4 Switzerland 0.955    83.8    16.3    13.4 69394.
# 5 Hong Kong, ~ 0.949    84.9    16.9    12.3 62985.
# 6 Iceland     0.949    83.0    19.1    12.8 54682.
```

# Remove unneeded rows

## Which rows are not needed?

```
head(HDI, n=2)
# # A tibble: 2 x 6
#   country      hdi life_exp school_exp school_mean gni_pc
#   <chr>      <dbl>   <dbl>     <dbl>     <dbl>   <dbl>
# 1 VERY HIGH H~ NA      NA      NA      NA      NA
# 2 Norway    0.957   82.4    18.1    12.9  66494.

HDI[HDI$country == "MEDIUM HUMAN DEVELOPMENT", ]
# # A tibble: 1 x 6
#   country      hdi life_exp school_exp school_mean gni_pc
#   <chr>      <dbl>   <dbl>     <dbl>     <dbl>   <dbl>
# 1 MEDIUM HUMAN~ NA      NA      NA      NA      NA
```

## Remove rows with no numeric data (e.g., in the hdi column).

```
HDI <- HDI[! is.na(HDI$hdi), ]

## Tidyverse alternative (NOT RUN)
## HDI <- dplyr::filter(HDI, ! is.na(HDI$hdi))
```

# Check the data **structure**

Things to check:

- Dimensions (rows by columns);
- Column names;
- Object “classes” (e.g., character vs. numeric).

```
str(HDI)
```

```
# tibble[,6] [189 x 6] (S3: tbl_df/tbl/data.frame)
```

```
# $ country      : chr [1:189] "Norway" "Ireland" "Switzerland" "Hong Ko
```

```
# $ hdi          : num [1:189] 0.957 0.955 0.955 0.949 0.949 0.947 0.945
```

```
# $ life_exp     : num [1:189] 82.4 82.3 83.8 84.9 83 ...
```

```
# $ school_exp   : num [1:189] 18.1 18.7 16.3 16.9 19.1 ...
```

```
# $ school_mean  : num [1:189] 12.9 12.7 13.4 12.3 12.8 ...
```

```
# $ gni_pc       : num [1:189] 66494 68371 69394 62985 54682 ...
```

# Check the data **summary**

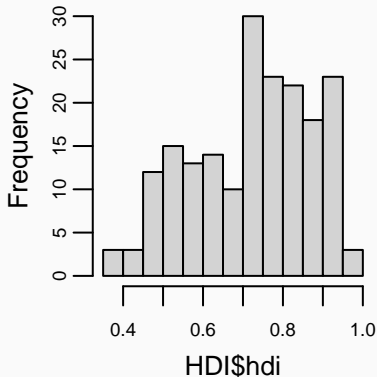
## Things to check:

- Descriptive statistics;
- Missing values (NA frequencies are reported for each variable).


```
summary(HDI)
#      country                hdi                life_exp
# Length:189             Min.   :0.3940             Min.   :53.28
# Class :character       1st Qu.:0.6020             1st Qu.:67.44
# Mode  :character       Median :0.7400             Median :74.05
#                               Mean  :0.7224             Mean  :72.71
#                               3rd Qu.:0.8290             3rd Qu.:77.91
#                               Max.   :0.9570             Max.   :84.86
#      school_exp      school_mean      gni_pc
# Min.   : 5.005      Min.   : 1.644      Min.   : 753.9
# 1st Qu.:11.431      1st Qu.: 6.437      1st Qu.: 4910.2
# Median :13.188      Median : 9.032      Median : 12707.4
# Mean   :13.325      Mean   : 8.728      Mean   : 20219.7
# 3rd Qu.:15.227      3rd Qu.:11.326      3rd Qu.: 29497.2
# Max.   :21.954      Max.   :14.152      Max.   :131031.6
```

## Run a few **hist**ograms?

```
hist(HDI$hdi,  
     main="")
```



```
hist(HDI$gni_pc,  
     freq=FALSE,  
     main="")  
lines(density(HDI$gni_pc),  
      col="blue")
```



## Step 1: Calculate the dimension indexes

The HDI is based on three “dimension” indexes:

**1 Health**

Based on life expectancy;

**2 Education**

Based on expected *and* mean years of schooling;

**3 Income**

Based on gross national income per capita.

## Calculating the dimension indexes

The dimension indexes are normalized using *modified min-max scales* with “goalpost” values:

$$\text{Dimension index} = \frac{\text{actual value} - \text{lower goalpost}}{\text{upper goalpost} - \text{lower goalpost}}$$

If the actual value falls between the goalposts, the dimension index value will fall between zero and one.

### The goalposts

Dimension	Indicator	Lower	Upper
Health	Life expectancy (yrs)	20	85
Education	Expected schooling (yrs)	0	18
	Mean schooling (yrs)	0	15
Income	GNI per capita (2017 PPP\$)	100	75,000

## Calculate the health index

Calculate the index using the “goalpost” formula, and assign it to the variable `health_index` in the HDI data frame.

```
HDI$health_index <- (HDI$life_exp - 20) / (85 - 20)
```

Check that the values are consistent with expectations.

```
summary(HDI$health_index)
#      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
#  0.5120  0.7298  0.8315  0.8109  0.8909  0.9978

summary(HDI$life_exp)
#      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
#   53.28   67.44   74.05   72.71   77.91   84.86
```



# Calculate the education index

There are two education indicators – the dimension index is the arithmetic mean of the two “goalposted” schooling indicators.

```
HDI$school_exp_index <- (HDI$school_exp - 0) / (18 - 0)
HDI$school_mean_index <- (HDI$school_mean - 0) / (15 - 0)
HDI$educ_index <- (HDI$school_exp_index + HDI$school_mean_index) / 2

## Check summary
summary(HDI[, c("school_exp_index", "school_mean_index", "educ_index")])
# school_exp_index school_mean_index educ_index
# Min. :0.2781 Min. :0.1096 Min. :0.2491
# 1st Qu.:0.6351 1st Qu.:0.4291 1st Qu.:0.5295
# Median :0.7327 Median :0.6021 Median :0.6823
# Mean :0.7403 Mean :0.5819 Mean :0.6611
# 3rd Qu.:0.8459 3rd Qu.:0.7551 3rd Qu.:0.7929
# Max. :1.2197 Max. :0.9434 Max. :1.0340
```

OOPS – what is the problem!

## “Enforce” the upper goalpost for school\_exp\_index

Index values must not exceed 1.0, even if the indicator exceeds the upper goalpost.

One way to fix this is to replace values of school\_exp\_index greater than 1.0 with values of exactly 1.0.

```
## Square brackets identify values to "reassign"
HDI[HDI$school_exp_index > 1, "school_exp_index"] <- 1

## Recalculate the dimension index
HDI$educ_index <- (HDI$school_exp_index + HDI$school_mean_index) / 2

## Check summary
summary(HDI[, c("school_exp_index", "school_mean_index", "educ_index")])
# school_exp_index school_mean_index educ_index
# Min.      :0.2781    Min.      :0.1096    Min.      :0.2491
# 1st Qu.:0.6351    1st Qu.:0.4291    1st Qu.:0.5295
# Median :0.7327    Median :0.6021    Median :0.6823
# Mean    :0.7366    Mean    :0.5819    Mean    :0.6592
# 3rd Qu.:0.8459    3rd Qu.:0.7551    3rd Qu.:0.7929
# Max.    :1.0000    Max.    :0.9434    Max.    :0.9433
```

# Calculate the income index

Calculate the income index, noting that it is based on the **natural logarithm** of GNI per capita.

The `log` function in R gives the natural logarithm by default.

```
HDI$income_index <- (log(HDI$gni_pc) - log(100)) /  
                    (log(75000) - log(100))  
summary(HDI[, c("gni_pc", "income_index")])  
#      gni_pc      income_index  
# Min.   : 753.9   Min.   :0.3051  
# 1st Qu.: 4910.2   1st Qu.:0.5882  
# Median : 12707.4   Median :0.7318  
# Mean   : 20219.7   Mean   :0.7152  
# 3rd Qu.: 29497.2   3rd Qu.:0.8590  
# Max.   :131031.6   Max.   :1.0843  
  
## Enforce the upper goalpost  
HDI[HDI$income_index > 1, "income_index"] <- 1  
summary(HDI$income_index)  
#      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.  
# 0.3051 0.5882 0.7318 0.7145 0.8590 1.0000
```

## Step 2: Combine the dimension indexes into “our” HDI

The HDI is the **geometric mean** of the three dimension indexes:

$$\text{HDI} = (\text{health} \times \text{education} \times \text{income})^{1/3}$$

(Raising to the one-third power is the same as taking a cube root.)

```
HDI$our_hdi <- (HDI$health_index * HDI$educ_index *  
               HDI$income_index) ^ (1/3)  
summary(HDI$our_hdi)  
#      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
# 0.3937  0.6019  0.7397  0.7225  0.8289  0.9570
```

# Compare “our” HDI with the UNDP’s

```
summary(HDI[, c("hdi", "our_hdi")])
#           hdi           our_hdi
# Min.      :0.3940    Min.      :0.3937
# 1st Qu.:0.6020    1st Qu.:0.6019
# Median :0.7400    Median :0.7397
# Mean     :0.7224    Mean     :0.7225
# 3rd Qu.:0.8290    3rd Qu.:0.8289
# Max.     :0.9570    Max.     :0.9570

## Differences between HDI values
HDI$hdi_diff <- HDI$hdi - HDI$our_hdi
summary(HDI$hdi_diff)
#           Min.      1st Qu.      Median      Mean      3rd Qu.
# -4.992e-04 -3.275e-04 -6.067e-05 -4.642e-05  2.258e-04
#           Max.
#  4.827e-04
```

Do we have a problem?

# The UNDP rounds HDI to the third decimal place!

What if we do the same?

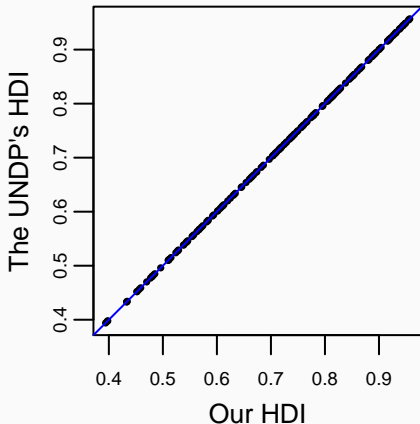
```
HDI$our_hdi <- round(HDI$our_hdi, digits=3)

## Recompare
HDI$hdi_diff <- HDI$hdi - HDI$our_hdi
summary(HDI$hdi_diff)
#      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
#         0         0         0         0         0         0
```

We have successfully wrangled the UNDP data in R and reproduced the HDI from the source data!

## Plotting our success

```
plot(x=HDI$our_hdi, y=HDI$hdi,  
     xlab="Our HDI",  
     ylab="The UNDP's HDI",  
     pch=16, cex=0.5)  
abline(a=0, b=1, col="blue") # x=y line
```



# Outline

On learning R

Some R basics

Data “wrangling” with R: the UNDP Human Development Index

**Model and visualize the data**

Self-study in R/RStudio with `swirl`

Wrapping up



## A regression model: income vs. non-income HDI

HDI was developed as an alternative to using national income (per capita) as a “proxy” measure of human development.

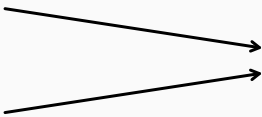
A concern beginning in the 1970s was that oil-exporting countries could have high income per capita but low human development.

### A (too?) simple model

Non-oil-rent national income  
(log, per capita)

Oil rents  
(log, per capita)

“Social” (non-income) HDI  
(health and education)




```
graph LR; A["Non-oil-rent national income (log, per capita)"] --> C["Social (non-income) HDI (health and education)"]; B["Oil rents (log, per capita)"] --> C;
```

## Calculate “non-income” HDI

Defining “non-income” HDI as the geometric mean of the health and education indexes, we can use our existing data to calculate it:

```
HDI$hdi_noninc <- sqrt(HDI$health_index * HDI$educ_index)
summary(HDI$hdi_noninc)
#      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
# 0.3897  0.6176  0.7585  0.7282  0.8382  0.9497

plot(x=HDI$income_index, y=HDI$hdi_noninc, cex=0.5)
abline(a=0, b=1, col="red", lwd=2) # x=y line
```



## Load data on oil rents

The HDI data set does not have oil rents. I downloaded some from the World Bank's *World Development Indicators*:

<https://databank.worldbank.org/source/world-development-indicators>

To save time, I put the data in an R data file, so that we only need to load the data (from the project folder).

```
load("oil_rents.RData")
summary(oil_rents)
#      isocode      oil_rents_pct
# Length:217      Min.   : 0.0000
# Class :character 1st Qu.: 0.0000
# Mode  :character Median : 0.0015
#                Mean   : 2.9722
#                3rd Qu.: 0.8795
#                Max.   :44.7903
#                NA's   :23
```

## Prepare to merge the oil data with the HDI data

Countries are identified in the oil-rent data using three-letter ISO country codes:

```
head(oil_rents$isocode)
# [1] "AFG" "ALB" "DZA" "ASM" "AND" "AGO"
```

To merge with the HDI data, we need the same three-letter country codes. Fortunately, the `countrycode` converts country names to country codes pretty well:

```
library(countrycode)
HDI$isocode <- countrycode(HDI$country,
                           origin="country.name",
                           destination="iso3c")
```

```
HDI[1:15, c("country", "isocode", "hdi")]
# # A tibble: 15 x 3
#   country          isocode    hdi
#   <chr>          <chr>    <dbl>
# 1 Norway         NOR      0.957
# 2 Ireland        IRL      0.955
# 3 Switzerland    CHE      0.955
# 4 Hong Kong, China (SAR) HKG      0.949
# 5 Iceland         ISL      0.949
# 6 Germany         DEU      0.947
# 7 Sweden          SWE      0.945
# 8 Australia       AUS      0.944
# 9 Netherlands     NLD      0.944
# 10 Denmark        DNK      0.94
# 11 Finland         FIN      0.938
# 12 Singapore       SGP      0.938
# 13 United Kingdom  GBR      0.932
# 14 Belgium         BEL      0.931
# 15 New Zealand     NZL      0.931
```

# Run the merge

Now we can merge the HDI and oil data.

```
dim(HDI) # dimensions before merging
# [1] 189  15

HDI <- merge(x=HDI, y=oil_rents,
             by="isocode",
             all.x=TRUE) # do not drop any "x" (HDI) data
```

And do a few quick checks on the merged data.

```
dim(HDI) # dimensions after merging (added one column)
# [1] 189  16

table(is.na(HDI$oil_rents_pct)) # "missing" oil-rent data
#
# FALSE  TRUE
#   181    8

rownames(HDI) <- HDI$isocode # this will be useful later
```

# Calculate oil-rent and non-oil-rent income

The oil-rent data is as a percentage of GNI.

```
HDI$gni_oil_pc      <- HDI$gni_pc * (HDI$oil_rents_pct / 100)
HDI$gni_nonoil_pc  <- HDI$gni_pc - HDI$gni_oil_pc
summary(HDI[, c("gni_oil_pc", "gni_nonoil_pc")])
#      gni_oil_pc      gni_nonoil_pc
# Min.      : 0.000   Min.      : 753.9
# 1st Qu.: 0.000   1st Qu.: 4857.8
# Median : 1.246   Median :12212.1
# Mean   : 739.236   Mean   :19120.9
# 3rd Qu.: 128.568   3rd Qu.:29368.7
# Max.   :25811.942   Max.   :88155.2
# NA's   :8         NA's   :8
```

Many countries have zero oil rents. Before taking logarithms, add one dollar to avoid undefined values.

```
HDI$gni_oil_pc_log    <- log(HDI$gni_oil_pc + 1)
HDI$gni_nonoil_pc_log <- log(HDI$gni_nonoil_pc)
```

## R's “formula” syntax for models

Models are specified in R using its “formula” syntax:

```
y ~ x1 + x2, data=df
```

- The response (dependent variable) is to the left of  $\sim$ ;
- Predictors (independent variables) are to the right of the  $\sim$ ;
- Predictors that enter “additively” are separated with plus signs (they are not literally added!);

Predictors with multiplicative interactions are separated by times signs (asterisks).

- A data argument allows to specify the data frame, so that you do not need to retype it for each variable in the formula.



## Running the linear model regression

The function for running a linear (OLS) regression in R is `lm`, which stands for “linear m.”

So we can run our model of the “non-income” part of HDI as a function of the (the log of) non-oil-rent GNI and oil rents as:

```
lm(hdi_noninc ~ gni_oil_pc_log + gni_nonoil_pc_log, data=HDI)
#
# Call:
# lm(formula = hdi_noninc ~ gni_oil_pc_log + gni_nonoil_pc_log,
#     data = HDI)
#
# Coefficients:
#           (Intercept)           gni_oil_pc_log  gni_nonoil_pc_log
#           -0.312802             -0.002105             0.112596
```

But the output is not very satisfying – just a bare printout of coefficients!

## Assign the model output!

The the `lm` output can be assigned as an object and summarized:

```
lm_out1 <- lm(hdi_noninc ~ gni_oil_pc_log + gni_nonoil_pc_log, data=HDI)
summary(lm_out1)
#
# Call:
# lm(formula = hdi_noninc ~ gni_oil_pc_log + gni_nonoil_pc_log,
#     data = HDI)
#
# Residuals:
#      Min       1Q   Median       3Q      Max
# -0.187955 -0.039392  0.008389  0.040400  0.147095
#
# Coefficients:
#              Estimate Std. Error t value Pr(>|t|)
# (Intercept)   -0.312802   0.037234  -8.401 1.37e-14 ***
# gni_oil_pc_log -0.002105   0.001541  -1.366   0.174
# gni_nonoil_pc_log 0.112596   0.004036  27.901 < 2e-16 ***
# ---
# Signif. codes:
#  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#
# Residual standard error: 0.06058 on 178 degrees of freedom
# (8 observations deleted due to missingness)
# Multiple R-squared:  0.8177, ^IAdjusted R-squared:  0.8156
# F-statistic: 399.2 on 2 and 178 DF, p-value: < 2.2e-16
```

# Extracting model estimates

The `lm` output contains model estimates that can be extracted:

```
names(lm_out1)
# [1] "coefficients" "residuals" "effects"
# [4] "rank" "fitted.values" "assign"
# [7] "qr" "df.residual" "na.action"
# [10] "xlevels" "call" "terms"
# [13] "model"





coef(lm_out1) # extract coefficients
# (Intercept) gni_oil_pc_log gni_nonoil_pc_log
# -0.31280185 -0.00210509 0.11259562

resid(lm_out1)[1:5] # extract and view the first five residuals
# AFG AGO ALB ARE ARG
# -0.02101448 -0.05647314 0.07106905 -0.05332631 0.06824457

fitted.values(lm_out1)[1:5] # same for the fitted values
# AFG AGO ALB ARE ARG
# 0.5551595 0.6188717 0.7490323 0.8991386 0.7951784
```

# Running “regression diagnostics”

```
par(mfrow=c(2, 2)) # set graphics to print "2 x 2"
plot(lm_out1)
```



# Outline

On learning R

Some R basics

Data “wrangling” with R: the UNDP Human Development Index

Model and visualize the data

Self-study in R/RStudio with `swirl`

Wrapping up

# What is `swirl`

## `swirl`

An R package that provides the infrastructure to run interactive self-study lessons, right in the R console.

(For more information: <https://swirlstats.com/>.)

## “Introduction to R Programming”

A foundational `swirl` course consisting of 15 short lessons (work out your own pace, but most take about 15–20 minutes each).

# Install and run swirl courses

## Install (once-off)

```
## Install the swirl package -- the "infrastructure"  
install.packages("swirl")  
  
## Install the R programming course -- the content  
swirl::install_course("R Programming")
```

## Run the course

```
## "Load" (attach) the swirl package  
library(swirl)  
  
## Run swirl  
swirl() # follow the prompts to choose a course, lessons
```

(Mostly do the courses in order, except you may want to skip the second one on “Workspace and Files” and come back to it later.)

# Outline

On learning R

Some R basics

Data “wrangling” with R: the UNDP Human Development Index

Model and visualize the data

Self-study in R/RStudio with `swirl`

Wrapping up



# Wrapping up

## Learning R

- Push ahead with the `swirl` lessons on “R Programming”;
- For further self-study, with a Tidyverse focus, try:  
Hadley Wickham and Garrett Grolemund, *R for Data Science* (O'Reilly Media, 2017) – available for free online at <https://r4ds.had.co.nz/>.

## For tomorrow

- Data *analysis* in R.